

Network Programming

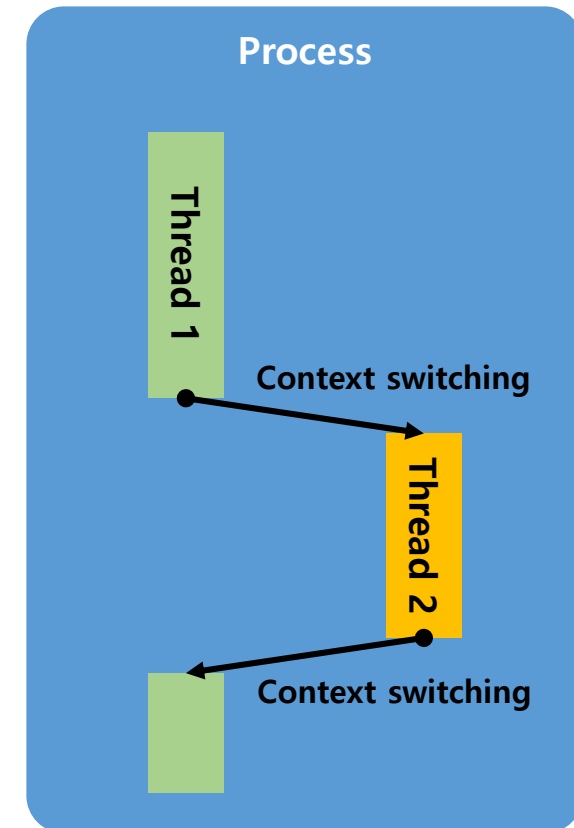
Thread Based Concurrent Programming

01

Thread

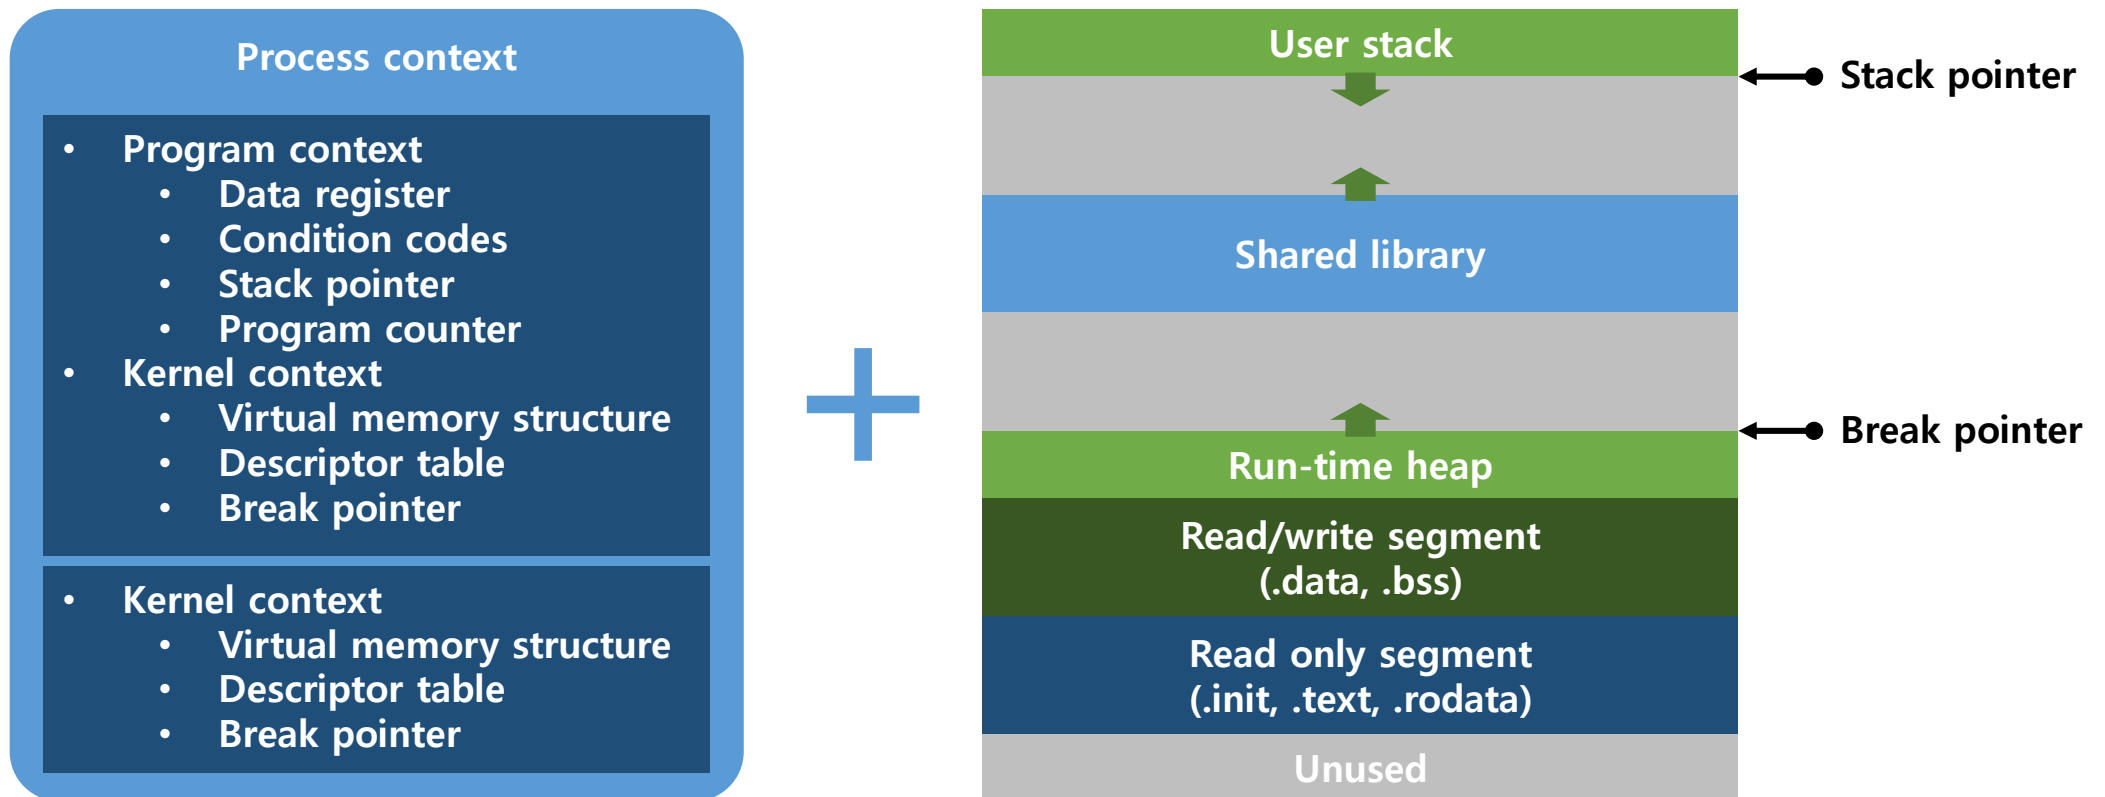
Thread

- The smallest sequence of programmed instructions
 - A part of a process
 - At least one thread exists in a process (**main thread**)
 - Threads in a process can share resources of their process
- Managed independently by a scheduler
 - Each thread has its control flow and address space
 - Context switching is performed between threads
 - Threads can run concurrently



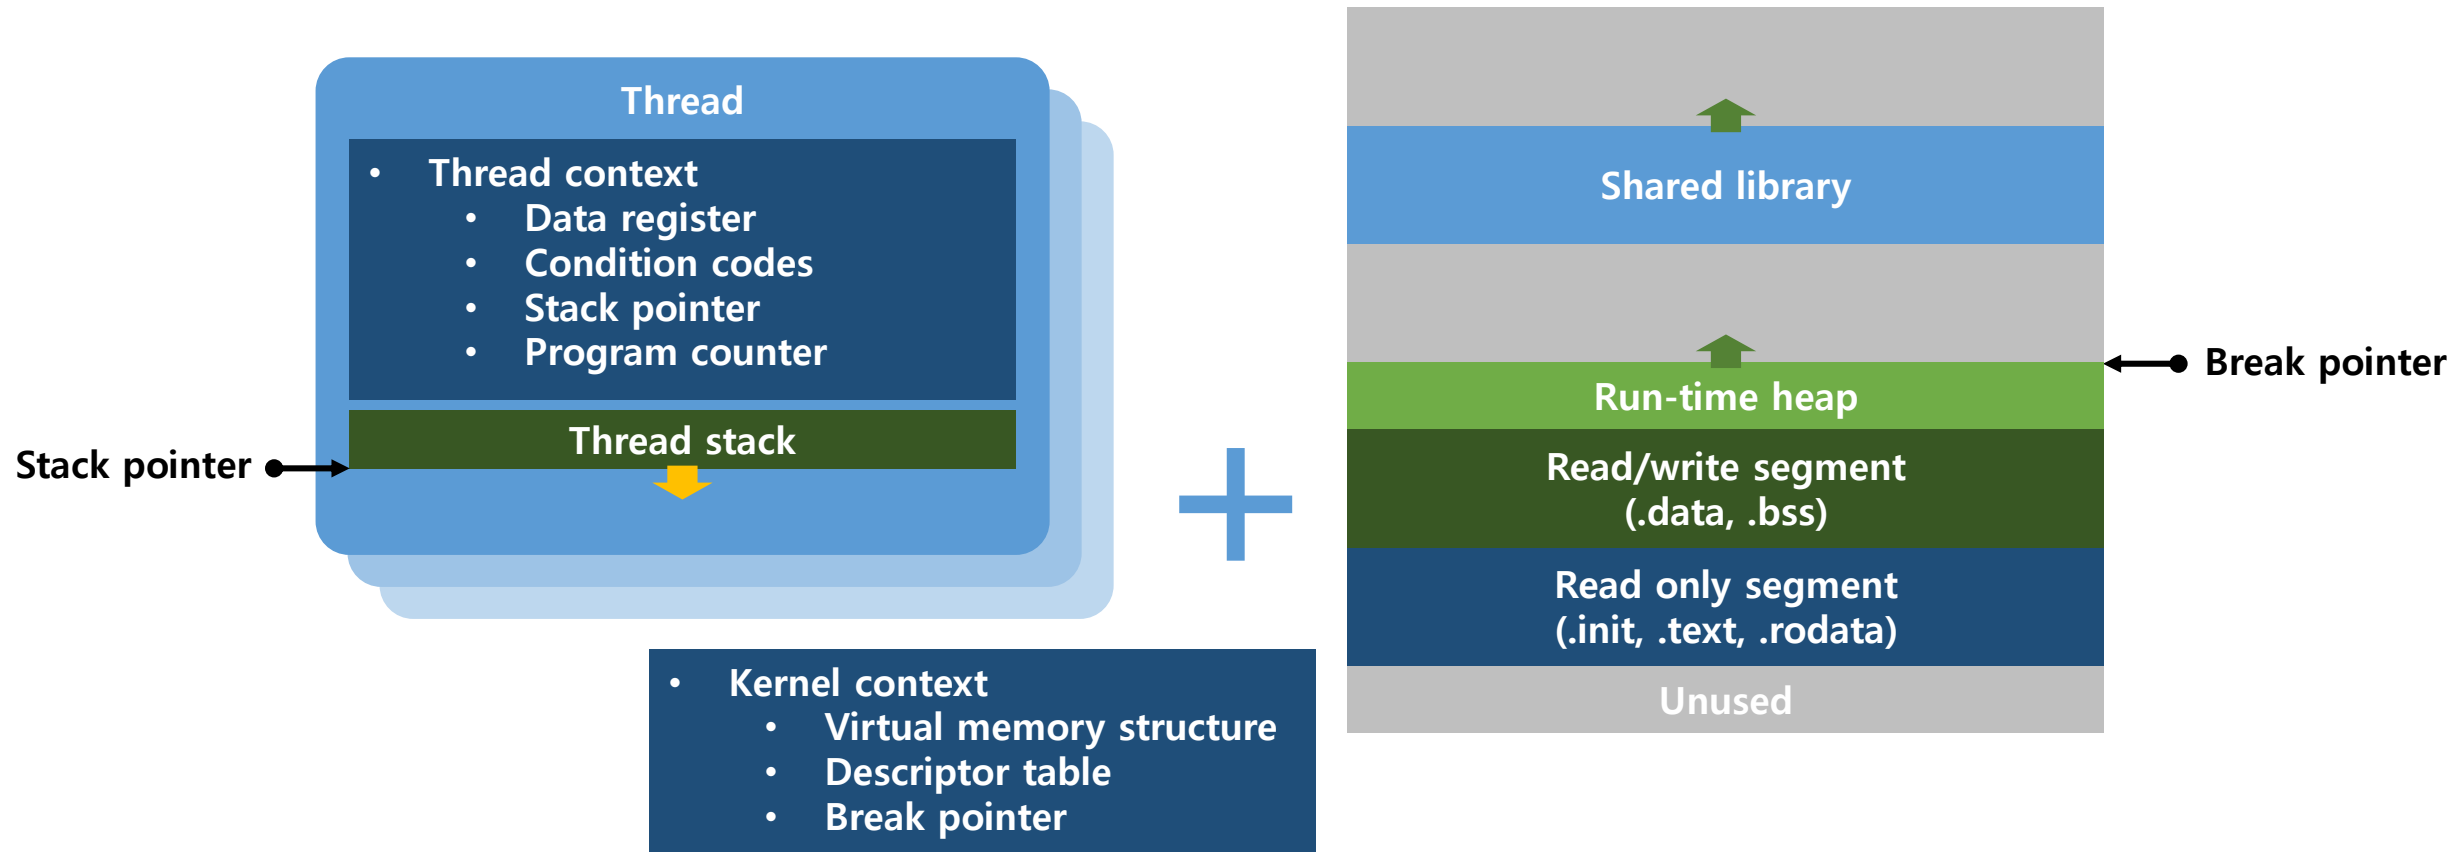
Traditional View of a Process

- Consist of **process context**, code, data, **stack**, and etc.



Alternative View of a Process

- Consist of **threads**, kernel context, code, data, and etc.




Alternative View of a Process


- Each thread has its own
 - Thread ID
 - Control flow
 - Stack and stack pointer
- Threads share
 - Kernel context of the process
 - Virtual memory address of the process

Logical View of Threads

- Processes form a **tree hierarchy**
- Threads associated with a process, form a **pool of peers**



Processes



Threads

Process and Thread

- Processes and threads are **similar** in
 - Each has its own ID
 - Each has its own control flow
 - Each has its own stack and stack pointer
 - Each can run concurrently
 - Each is context switched
- Processes and threads are **different** in
 - Threads can **share** resource of the process **easily**,
 - Processes can share resources too via IPC (inter process communication), but it's hard
 - A thread are **much more lighter** than a process
 - **Faster** to create and reap

Posix Thread

- Posix
 - Portable operating system interface
 - Family of **standards** specified by the IEEE computer society
 - Maintaining **compatibility** between operating systems based on **UNIX**
 - Defines the API for operating systems based on UNIX
- Pthread
 - **Standard thread interface** for UNIX systems
 - About 60 functions are supported for C language

Posix Thread

- Posix
 - Portable operating system interface
 - Family of **standards** specified by the IEEE computer society
 - Maintaining **compatibility** between operating systems based on **UNIX**
 - Defines the API for operating systems based on UNIX
- Pthread
 - **Standard thread interface** for UNIX systems
 - About 60 functions are supported for C language

Posix Thread Creation

- Posix thread create function

- `int pthread_create(pthread_t* threadID, const pthread_attr_t* attribute, void* startRoutine, void* argumentPointer);`

- **Create** a new thread and **run a routine** with it
 - **threadID**: the ID of the new thread will be stored in here
 - **attribute**: attribute structure for the thread, usually NULL
 - **startRoutine**: a routine to run with the thread
 - **argumentPointer**: a pointer of arguments for the thread
 - Return **0** on **success**, **non-zero** on **error**

Posix Thread Joining

- Posix thread join function
 - `int pthread_join(pthread_t* threadID, void** returnValuePointer);`
 - **Block** the caller, and **return** when specified **thread terminates**
 - `threadID`: the ID of a thread to wait for
 - `returnValuePointer`: a return value of the thread will be stored in here
 - Return **0** on **success**, **non-zero** on **error**

Posix Thread Creation and Joining

```
...
#include <pthread.h>
// Add an option -lpthread to compile

void* threadRoutine(void* argumentPointer){
    pthread_t id = pthread_self();
    // Return the ID of the calling thread
    printf("Hello, this is thread %d.\n", id);
    return NULL;
    // Thread should return something for join function
}

int main(){
    pthread_t threadID;

    pthread_create(&threadID, NULL, threadRoutine, NULL);
    pthread_join(threadID, NULL);

    return 0;
}
```

Passing Arguments to a Thread

```
...
void* threadRoutine(void* argumentPointer){
    argument = (argumentType) *argumentPointer;
    // Cast the argumentPointer into a proper type
    ...
    // You can use the argument in here
}

int main(){
    pthread_t threadID;

    argumentType argument;
    // An argument of a type

    pthread_create(&threadID, NULL, threadRoutine, (void*) &argument);
    // Should pass the argument as a void pointer type
    // Of course you can just cast an integer into a void pointer type to pass
    ...
}
```

Passing a Return Value to Join Function

```
...
void* threadRoutine(void* argumentPointer){
    ...
    int returnValue = 2016;
    return (void*)returnValue;
    // Should return a value as void pointer type
}

int main(){
    ...
    void* returnValuePointer;
    // Should be void pointer type
    pthread_join(threadID, &returnValuePointer);

    int returnValue = (int)returnValuePointer;
    // Cast the returnValuePointer into a proper type
    // Of course you can allocate a new return value than deallocate it
    ...
}
```

Posix Thread Detach

- Posix thread detach function
 - `int pthread_detach(pthread_t threadID);`
 - Detach the thread
 - **Don't wait** the thread with **join function** detached threads
 - When detached thread terminates, system reaps it **automatically**
 - threadID: the ID of a thread to detach
 - Return **0** on **success**, **non-zero** on **error**

Posix Thread Detach

```
...  
# include <pthread.h>  
  
...  
  
int main(){  
    pthread_t threadID;  
  
    pthread_create(&threadID, NULL, threadRoutine, NULL);  
    pthread_detach(threadID, NULL)  
  
    ...  
    // Don't call a join function for the threadID  
  
    return 0;  
}
```

02

Memory Sharing

Memory Sharing

- Threads can **share** the virtual memory of their **process**
- Threads **can not** share each own **stack**
 - Note that the **main function** is a thread too
 - Hence, threads can share **global variables** of their process
- Threads can share **static** variables in each own **stack**

Memory Sharing Example

```
int globalVariable = 2016;

void* threadRoutine(void* argumentPointer){
    static int staticVariable = 1;
    int localVariable = pthread_self();
    ...
}

int main(){
    ...

    pthread_create(threadIDs[0], NULL, threadRoutine, NULL);
    pthread_create(threadIDs[1], NULL, threadRoutine, NULL);
    ...
    int localVariable = 1027;
    int* localPointer = (int*)malloc(...);
    ...
}
```

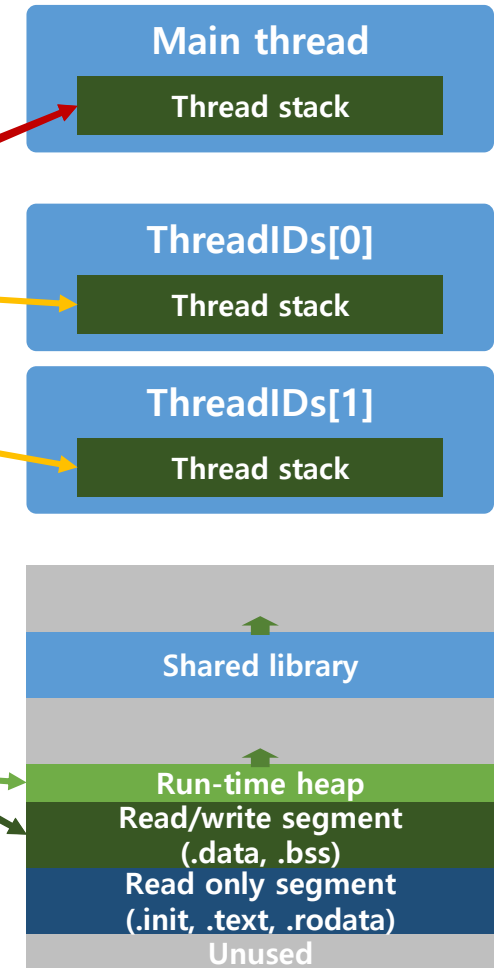
Memory Sharing Example

```
int globalVariable = 2016;

void* threadRoutine(void* argumentPointer){
    static int staticVariable = 1;
    int localVariable = pthread_self();
    ...
}

int main(){
    ...

    pthread_create(threadIDs[0], NULL, threadRoutine, NULL);
    pthread_create(threadIDs[1], NULL, threadRoutine, NULL);
    ...
    int localVariable = 1027;
    int* localPointer = (int*)malloc(...);
    ...
}
```



Memory Sharing Example

- Accessible table

	Main thread	threadIDs[0]	threadIDs[1]
globalVariable	O	O	O
staticVariable	X	O	O
localVariable (in threadIDs[0])	X	O	X
localVariable (in threadIDs[1])	X	X	O
localVariable (in main thread)	O	X	X
localPointer	O	X	X

Unsafe Region

- Memory sharing makes the application **efficient**
 - Much less redundant than using fork()
 - Faster to create and reap than processes
- However, it can be very **dangerous**
 - Memory sharing can make **unsafe region** in the application

Unsafe Region Example

```
int sum = 0;

void* threadRoutine(void* argumentPointer){
    for(int i = 0; i < 1000000; i++)
        sum++;
    ...
}

int main(){
    ...

    pthread_create(threadIDs[0], NULL, threadRoutine, NULL);
    pthread_create(threadIDs[1], NULL, threadRoutine, NULL);

    pthread_join(threadIDs[0], NULL);
    pthread_join(threadIDs[1], NULL);

    printf("sum: %d\n", sum);
    ...
}
```


Unsafe Region Example

- The result of sum should be **2,000,000**
 - But it's **not** (in a high probability)
- One line of C source code is **not one instruction** for CPU
 - In pseudo assembly code, `sum++` is
 - `load sum eax`
 - `increment eax`
 - `store sum eax`

Unsafe Region Example

- If the threads run like below, the result must be 2,000,000

ThreadIDs[0]

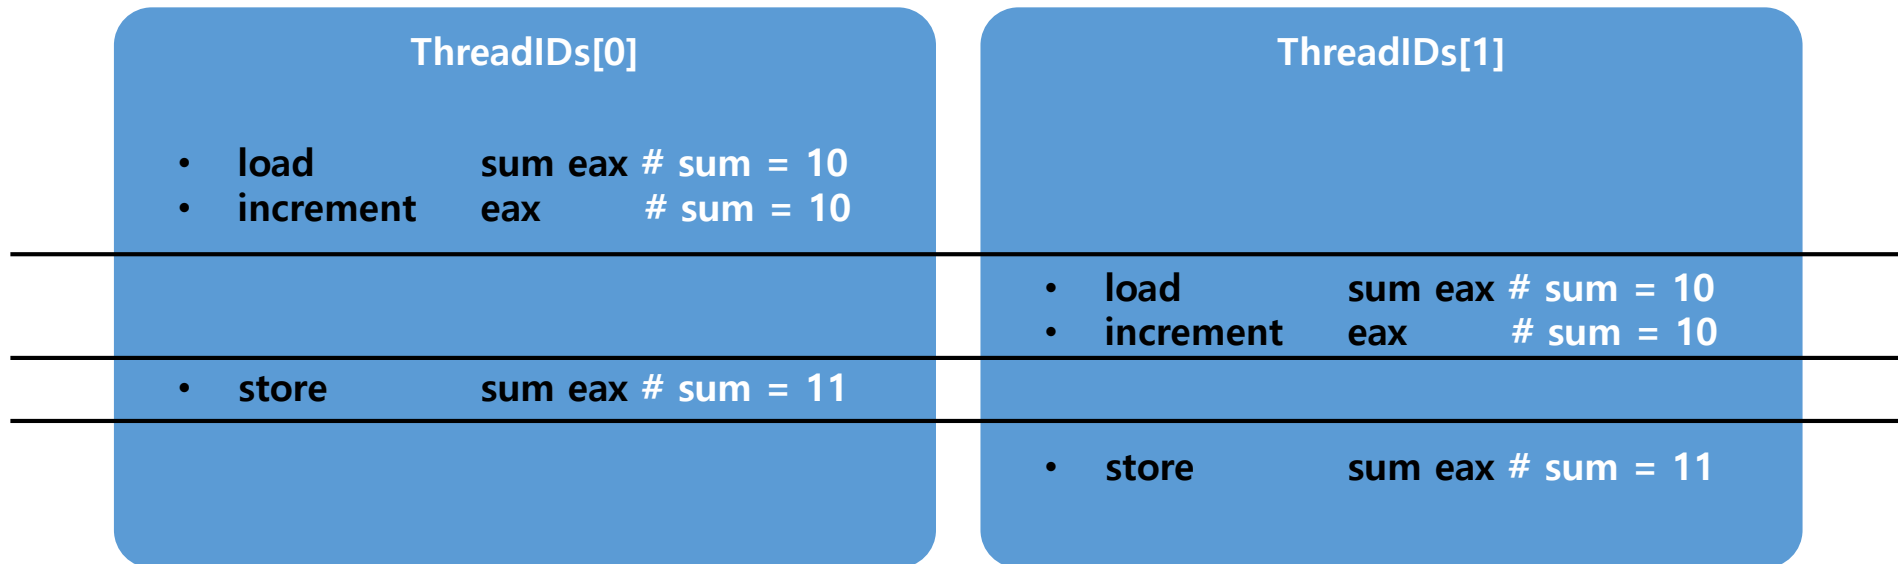
- load sum eax # sum = 10
- increment eax # sum = 10
- store sum eax # sum = 11

ThreadIDs[1]

- load sum eax # sum = 11
- increment eax # sum = 11
- store sum eax # sum = 12

Unsafe Region Example

- However, threads don't run like above in a high probability
 - For example like below



Unsafe Region Example

```
int sum = 0;

void* threadRoutine(void* argumentPointer){
    for(int i = 0; i < 1000000; i++)
        sum++; // This area is unsafe region
    ...
}

int main(){
    ...

    pthread_create(threadIDs[0], NULL, threadRoutine, NULL);
    pthread_create(threadIDs[1], NULL, threadRoutine, NULL);

    pthread_join(threadIDs[0], NULL);
    pthread_join(threadIDs[1], NULL);

    printf("sum: %d\n", sum);
    ...
}
```

Synchronizing

- Synchronizing enforce **concurrent** flow to be **sequential**
 - Synchronizing unsafe region prevent **unintended** context switching
- Synchronizing consist of two steps
 - Wait
 - Check some conditions **before** run a synchronized region
 - If the **conditions** are not satisfied, wait for a **wake up signal**
 - When the **signal** is received, it wakes up and test some **conditions** again
 - Signaling
 - Send a **wake up signal** when **exiting** the synchronized region
- Two ways of multi-thread synchronizing
 - Mutual exclusion
 - Semaphore

Mutex

- Mutual exclusion
 - Set **critical section** to synchronize an unsafe region
 - **Only one** thread can run critical section
- Declaring a mutex
 - Use **PTHREAD_MUTEX_INITIALIZER** macro for **static** mutex
 - Use **pthread_mutex_init()** function for **dynamic** mutex
- Destroying a mutex
 - Use **pthread_mutex_destroy()** function
 - If the mutex is **locked**, return **EBUSY**
 - Hence, the mutex **should be unlocked** before being destroyed

Mutex

```
#include <pthread.h>

pthread_mutex_t counter_mutex = PTHREAD_MUTEX_INITIALIZER;
// Static mutex declaration

void* threadRoutine(void* argumentPointer){
    ...
    pthread_mutex_lock(&counter_mutex);
    ...
    // Critical section with a mutex
    pthread_mutex_unlock(&counter_mutex);
    ...
    return NULL;
}
```

Mutex

```
int main(){
    ...

    pthread_create(threadIDs[0], NULL, threadRoutine, NULL);
    pthread_create(threadIDs[1], NULL, threadRoutine, NULL);

    pthread_join(threadIDs[0], NULL);
    pthread_join(threadIDs[1], NULL);

    pthread_mutex_unlock(&counter_mutex);
    // Should unlock before destroy
    pthread_mutex_destroy(&counter_mutex);

    return 0;
}
```


Semaphore

- Let **only limited threads** run a synchronized region
- Dijkstra's P and V operations
 - P(s) operation
 - Proberen ("test" in Dutch)
 - It works like [while(**s==0**) wait(); s--;]
 - V(s) operation
 - Verhogen ("**increment**" in Dutch)
 - It works like [**s++**;
 - P and V operations are atomic operations

Semaphore

```
...  
# include <semaphore.h>  
...  
  
sem_t semaphore;  
int limit = 1;  
  
sem_init(&semaphore, limit, 0);  
// Initialize a semaphore  
// limit: max number of threads for the synchronized region  
// 0: initial value for the semaphore  
  
sem_wait(&semaphore);  
// P operation  
...  
// Synchronized region with a semaphore  
sem_post(&semaphore);  
// V operation  
  
sem_destroy(&semaphore);
```

Problems of Multi-thread Programming

- Race condition
 - Occur because **no one knows** which thread run first
 - Some **results can be changed** by the **order** of thread execution
- Unintended context switching
 - Can prevent by using **synchronizing**
- Deadlock
 - Occur because of **synchronizing**
 - Wrongly designed synchronizing makes **every thread wait** for the wake up signal
 - The whole process is **halt**